

Service Level Objectives

Online services want to ensure customers have a positive experience in their product. One way they do this is by defining **Service Level Objectives**, or **SLOs**.

An SLO is like a guarantee made to customers. By defining an SLO, an organization is promising to maintain a defined level of service, often related to the performance and availability of their systems. Here are some examples:

- The website landing/home page will take <5 seconds to load 99% of the time over a 7 day period
- Checkout service will operate without error 99% of the time each 30 days
- The landing/home page will be able to successfully process at least 1,000 requests per second, 99% of the time, measured in 90-day increments
- Average response time in the cart service is <300 milliseconds, 98% of the time, measured in 7-day increments

Notice each includes the same components:

- **Scope:** The specific area an SLO relates to, such as “checkout service”, or “landing page”. This is an area or function that impacts user experience.
- **Target Value:** A measurable threshold for performance, like “1,000 requests per second” or “less than 5 seconds”. Data used to measure an SLO is referred to as a **Service Level Indicator (SLI)**, because it *indicates* whether an SLO is met.
- **Target Rate:** A percentage of the time performance will meet the target value.
- **Time Window:** Period for which data is evaluated, such as “over 7 days”.

Notice also that each SLO is centered around user experience. SLOs should capture performance and availability levels that, even if *barely* met, would keep the average user satisfied. In the simplest terms:

- **Service *meets* SLO targets → Happy users**
- **Service *misses* SLO targets → Frustrated users**

For this reason, SLOs monitor the performance and availability most likely to impact a user’s experience: like load times, error rates, and functionality working in the manner users need and expect.

Service Level Indicators, Objectives, and Agreements

An SLO should be focused on user experience. Think about what makes a positive user experience in a website, app, or online service:

- Load times are reasonable
- Functionality works as expected
- Users are able to complete whatever task or action they visited the service for, for example:
 - On an ecommerce site, you can complete an order
 - On a digital calendar, you can create, edit, and delete calendar events without error

For example, a typical user won't care about the precise integrity of your internal systems. But they *will* care if those systems negatively impact their experience with long load times, frequent errors, and inconsistent functionality.

When considering SLOs, focus on users' goals and experiences. Ask yourself questions like:

- How are your users interacting with your application?
- What is their journey through the application?
- Which parts of your infrastructure do these journeys interact with?
- What are they expecting from your systems? What are they hoping to accomplish?
- What would prevent them from accomplishing their goals?

After identifying what is necessary to keep users content, the next step is locating data to measure each SLO.

Service level indicators (SLIs)

Data used to measure an SLO is called a **Service level indicator (SLI)**, because it *indicates* whether the SLO is met.

For example, the SLO *"The website landing/home page will take <5 seconds to load 99% of the time over a 7 day period"* refers to website load times. To determine whether this SLO is met, you would need latency data capturing average landing page load times for the past 7 days. This latency data would be the SLI for this SLO.

Because SLOs aim to quantify user experience, their associated SLIs are values that correspond to the quality of that user experience. This includes things like:

- The number of requests to an endpoint that complete successfully
- The number of requests to an endpoint that complete within 500ms
- Average load times in specific areas and pages users need

SLIs are often formatted as percentages, representing the rate of “good” events out of all valid events. For example, a valid event could be a user request to an endpoint, regardless of the request’s success or failure. A “good” event would be a successful `200 OK` request to that endpoint. The SLI would be the percentage of all valid endpoint requests that were `200 OK` successful.

$$\text{SLI} : \frac{\text{good events}}{\text{valid events}} \times 100\%$$

Service level agreements (SLAs)

Service level agreements, or **SLAs**, address expectations, impacts, and consequences if agreed-upon reliability is not met. This is a contract that outlines what happens if a service meets—or misses—SLO targets. Consequences for missing SLOs are often business decisions or financial repercussions, such as:

- The customer expects a given service to have a 0.05% maximum error rate daily, or they’ll receive a rebate.
- The customer expects only 10% of monthly requests to take longer than 500ms to complete, or they’ll be reimbursed for the compute overage.

These stipulations are often related to business and financial relationships between the user and service. As such, SLAs are generally outside of the purview of Engineers. Instead, they’re often written by other company stakeholders before specific SLOs are quantified.

Engineering teams can derive their SLOs from the SLAs defined by other parts of the company, but SLO targets should generally be more strict than SLA targets.

Time window

In an SLO, the time window refers to the period for which the SLI will be evaluated to determine if it meets SLO targets.

It's best practice to define a rolling time window—such as `the past 30 days`—as opposed to a fixed window, like `the month of August`. This is because a user's perception of a service's reliability is heavily influenced by their recent experience with that service. The average user will not reset their expectations when a new month begins.

Typical SLO time windows are 7, 30, and 90 days.

Target rates & reliability

You may have noticed SLO target rates are often ~99%, as seen in examples from previous lessons. But if SLOs are meant to guarantee positive user experiences, why aren't they aiming for 100%?

In reality, 100% reliability is an impractical target that goes against best practices for many reasons:

- **It's difficult:** Striving for perfection is hard, especially over longer periods of time. It puts undue stress and expectations on engineering teams.
- **Users don't need it:** The difference between 99% and 100% is negligible—even unnoticeable—for most users.
- **Teams must balance time:** Engineers can't spend all their time maintaining perfect reliability and availability. Some of their time must be used to deliver value through developing new features, products, etc.

Aiming for less than 100% reliability means a small amount of undesirable performance/events are allowable. For example, if an SLO has a target of 99% uptime, outages could still occur $\geq 1\%$ of the time without breaching the SLO. This amount of allowable "bad" events is called an **error budget**.

Error budgets

An error budget quantifies an acceptable amount of unreliability. In other words, the amount of "bad" events that can happen without breaching an SLO. They are formatted as 100% minus the SLO target. Consider the following:

```
Average cart service response time is <300 milliseconds, 98% of the time,  
measured in 7-day increments.
```

This SLO targets a response time of <300ms 98% of the time in a rolling 7-day window. This means $\geq 2\%$ of requests may have *longer* response times, without breaching the SLO. For every 100,000 requests, 2,000 are “allowed” to fail in a 7-day period, because the SLI’s total rate of will still be $\geq 98\%$.

How teams use error budgets

Like a financial budget, an error budget is meant to be spent. When a team knows how much budget is remaining in their SLOs, they can allocate time and resources accordingly.

$$\text{Error Budget} = 100\% - \text{SLO Target}$$

When an error budget is nearly exhausted ($\leq 0\%$ remaining), teams prioritize reliability to ensure they don’t breach SLOs. This means focusing on lower-risk tasks unlikely to consume budget, such as:

- Prioritizing postmortem items
- Automating deployment pipelines
- Improving monitoring and observability
- Freezing feature releases

When an error budget is readily available ($> 0\%$ remaining), teams prioritize velocity and higher-risk tasks, because they have ‘room’ to encounter errors and issues without breaching SLOs. This includes things like:

- Releasing new features
- Making system changes
- Trying risky (but valuable) experiments

Burn rates

It’s also helpful to understand when error budgets are being spent at a higher rate than usual. This is tracked by something called a burn rate. A **burn rate** is a unitless number indicating how fast an error budget is being “spent” relative to its SLO target.

For example, if an SLO with a 30-day window is on track to consume its budget in 30 days, it would have a burn rate of **1**. If it was on track to consume its budget in 15 days it would have a burn rate of **2**, because it is “spending” its error budget at 2x the intended rate.

By tracking burn rates and identifying when rates spike or increase, you can resolve issues and return the rate to baseline before your SLO's error budget is consumed.

SLOs in Practice

2. In your browser, log in to app.datadoghq.com
3. In the main menu, hover over **APM** and click **Traces** to navigate to the [Traces Explorer](#).
4. Locate the facet panel on the left. Expand options for **Env** and **Service**.

Service Level Objectives—or SLOs—define the level of service you strive to provide users. To identify SLOs, you must understand the user experience your application or product offers. This includes things like:

- What do users achieve in your application?
- What are the most critical pathways, tasks, or actions?
- What would prevent them from having a good—or even functional—experience?

5. Navigate to [APM > Services > Service Catalog](#). This is a list of all services associated with your account. Locate the env dropdown menu above the list. This is where you select your environment
6. Hover over the service you'd like to inspect in the list entry to reveal a Full Page button. Click Full Page to open the service overview page
7. Locate the Resources list near the bottom of the page.
8. In the Resources list on the service overview page, click your desired `route/endpoint` to open its details.
9. Under Resource Summary, locate the Latency graph. Click the View full screen button (four outward arrows) above the graph to enlarge. Notice there are multiple metrics in the legend below the graph. Each metric name is preceded by p + a numerical value (for example, p50, p75, etc). This refers to a distribution percentile. The ideal percentile would be metric p99 this means that 99% of requests have latencies lower than this value.

Creating a Monitor for SLO's (in practice)

Some screenshots were taken from a Datadog lab. Just replace with your `env`, `service`, `route/endpoint`.

Datadog offers two approaches to measuring SLOs:

1. **By Count:** References metric(s) tracked in Datadog to measure the ratio of good events to total events. The SLO target is the percentage of how many events were "good" in a given time period.

- These are also called **Metric-Based SLOs**.

2. **By Monitor Uptime:** References one or more existing monitors. The SLO target is the percentage of time monitors are in an `OK` state. This is most useful for time-based data.

- These are also known as **Monitor-Based SLOs**.

Your SLO for homepage load times will be measured by Monitor uptime, because latency is a time-based metric. It will look like this:

"Over the past 7 days, 99% of the time the p99 latency of a home page request should be less than 5 seconds."

That is, for 99% of the past 7 days, 99% of home page requests should experience less than 5 seconds of latency.

Our SLO will be Monitor-based, you'll need to create a new Monitor for homepage latency. The Monitor will then be used to construct your SLO.

1. In Datadog, navigate to **Monitors > New Monitor**.

2. Select **APM** from the list on the left. This will open a form.

3. Under **Select monitor scope**, select **APM Metrics**.

- In the **Primary tags** field, select `env:{yourenv}`.
- For **Service**, select `service-yourservice`.
- For **Resource**, select `route/endpoint`.

- For **Evaluate the query over the last**, leave the default **5 minutes** .

1 Select monitor scope

APM Metrics

Trace Analytics

Watchdog

Primary tags:

env:ruby-shop

Service:

store-frontend

Resource:

spree::homecontroller_index

Evaluate the query over the last

5 minutes

4. Under **Set alert conditions**, select **Threshold Alert** .

- Set the alert criteria to **Trigger when** **p99 latency** **is** **above** **the threshold**.
- For the **Alert threshold** enter **5** .
- Leave the optional **Warning threshold** blank.

2 Set alert conditions

Threshold Alert

Anomaly Alert

An alert is triggered whenever a metric crosses a threshold. ?

Trigger when **p99 latency** is **above** the threshold

Alert threshold: >

5

Warning threshold: >

Optional

> Advanced options

5. In the **Notify your team** section you will see a text editor.

- After several moments it should update to an status more than likely it'll be **OK**.

Monitors > Status [Mute] [Escalate] [Settings]

OK Monitor status since 1 minute ago (13 Jul, 11:32:17 America/Los_Angeles)

Resource spree::homecontroller_index has an abnormal change in throughput on env:ruby-shop

Properties

APM Monitor
ID: 125233457
Created at Jul 13, 2023, 11:33 am
by dkvpc32hab

QUERY

```
avg(last_5m):p99:trace.rack.request{env:ruby-shop, resource_name:spree::homecontroller_index, service:store-frontend} > 5
```

TAGS

- env:ruby-shop
- service:store-frontend
- resource_name:spree::ho...

RECIPIENT

No recipients defined

MESSAGE

Monitor for the SLO "99% of the time the p99 latency of a home page request should be less than 5 seconds."

TEAMS

PRIORITY

Not Defined

- On the detail page for your new Monitor, notice the metric query in the **Query** section.

Monitors > Status [Mute] [Escalate] [Settings]

OK Monitor status since 1 minute ago (13 Jul, 11:32:17 America/Los_Angeles)

Resource spree::homecontroller_index has an abnormal change in throughput on env:ruby-shop

Properties

APM Monitor
ID: 125233457
Created at Jul 13, 2023, 11:33 am
by dkvpc32hab

QUERY

```
avg(last_5m):p99:trace.rack.request{env:ruby-shop, resource_name:spree::homecontroller_index, service:store-frontend} > 5
```

TAGS

- env:ruby-shop
- service:store-frontend
- resource_name:spree::ho...

RECIPIENT

No recipients defined

MESSAGE

Monitor for the SLO "99% of the time the p99 latency of a home page request should be less than 5 seconds."

TEAMS

PRIORITY

Not Defined

- The `trace.rack.request` metric is scoped to the respective `env`, `service`, and `resource_name`.

Note: For more information on distribution trace metrics, review the [DDSketch-based Metrics in APM](#) documentation.

8. Notice that **Tags** have also been added for `env` , `service` , and `resource_name` .

The screenshot shows the Datadog 'Monitors > Status' page. At the top, a green status bar indicates 'OK' and 'Monitor status since 1 minute ago (13 Jul, 11:32:17 America/Los_Angeles)'. Below this, a green alert message states: 'Resource spree::homecontroller_index has an abnormal change in throughput on env:ruby-shop'. The main content area is titled 'Properties' and contains details for an 'APM Monitor' with ID 125233457, created on Jul 13, 2023, at 11:33 am by user dkvpc32hab. A 'TAGS' section lists three tags: 'env:ruby-shop', 'service:store-frontend', and 'resource_name:spree::ho...'. Below the tags is a 'TEAMS' section and a 'PRIORITY' section set to 'Not Defined'. To the right of the properties, there are three sections: 'QUERY' showing a SQL query for p99 latency, 'RECIPIENT' stating 'No recipients defined', and 'MESSAGE' describing the SLO: 'Monitor for the SLO "99% of the time the p99 latency of a home page request should be less than 5 seconds."'.

Monitors > Status

OK Monitor status since 1 minute ago (13 Jul, 11:32:17 America/Los_Angeles)

Resource spree::homecontroller_index has an abnormal change in throughput on env:ruby-shop

▼ Properties

APM Monitor
ID: 125233457
Created at Jul 13, 2023, 11:33 am
by dkvpc32hab

TAGS

- env:ruby-shop
- service:store-frontend
- resource_name:spree::ho...

TEAMS

PRIORITY Not Defined

QUERY

```
avg(last_5m):p99:trace.rack.request{env:ruby-shop, resource_name:spree::homecontroller_index, service:store-frontend} > 5
```

RECIPIENT No recipients defined

MESSAGE Monitor for the SLO "99% of the time the p99 latency of a home page request should be less than 5 seconds."

- You can use these tags to search the monitors in the **Manage Monitors** list. They also link the monitor to corresponding APM Service and Resource pages.

With your monitor in place, you can now use it to create your Monitor-based SLO for your page latency issue.

Creating a Monitor based SLO (in practice)

1. In Datadog, navigate to **Service Mgmt > Services > SLOs**.

2. Click **New SLO** in the upper-right. This will open a form.

SLOs > Create SLO

1 Define your SLO measurement

Select how to measure your SLO

✓ By Count

Measures the ratio of good events to total events.

🕒 By Monitor Uptime

Measures the uptime of your monitors.

Define your SLI

Count-based SLOs are calculated using a percentage of good events out of total events (Good Events + Total Events × 100). [Need help?](#) [See our docs](#)

Good Events

a (select metric) from (everywhere) sum by (everything) </>

as count

+ Add Query + Add Formula

Total Events

a (select metric) from (everywhere) sum by (everything) </>

as count

+ Add Query + Add Formula

2 Set your target & time window

Evaluate over a rolling time window of 7 Days with a target of 99.9 %

OPTIONAL

Warning threshold (99.95) %

3 Add name and tags

Name

Description

Edit Preview H B I S | 🔗 " " </> 📄 ☰ ☷ ☑️ 🖨️ ⌵

Write an additional description

Tags

Add tags

Team Tags

Add teams

Cancel

Create & Set Alert

Create

Preview

Select metrics to see your SLO preview

3. Under **Define your SLO measurement**, select **By Monitor Uptime**. This will reveal a **Select monitors** dropdown menu.

4. In the **Select monitors** dropdown, select the monitor you created above: " route/endpoint has an abnormal change in throughput on env : {yourenv} "

1

Define your SLO measurement

Select how to measure your SLO

By Count

Measures the ratio of good events to total events.

✓

By Monitor Uptime

Measures the uptime of your monitors.

Select monitors

Add up to 20 monitors to use for SLO calculation

ADDED MONITORS

Resource spree::homecontroller_index has an abnormal change in throughput on env:ruby-shop

Search and add a monitor

5. Under **Set your target & time window**, update the statement to read: **Evaluate over a rolling time window of 7 days with a target of 99%**.

- Leave the optional **Warning threshold** blank.

2 Set your target & time window

Evaluate over a rolling time window of with a target of %

OPTIONAL

☐ Warning threshold %

6. Under **Add names and tags**, enter the following **Name**: `Page P99 Latency`


Under **Tags**, select each of the following:

- `env :`
- `service:yourservice`
- `route/endpoint`

Add **Team Tags** as-needed.

Click **Create**. You will navigate to a detail panel for your new SLO.

The **History** section is might be empty, but this is expected if this is a new `route/endpoint`.



SLO replay: Monitor-based or "By Monitor Uptime" SLOs also have a feature called **SLO Replay**. SLO Replay backfills SLO statuses with historical data pulled from underlying monitors' metrics and queries. This means if you create or update a Monitor-based SLO it will be populated with relevant historical data when available.

Creating a Metric based SLO (in practice)

As determined earlier, the second Storedog SLO is: `"Over the past 30 days, 99% of requests to the cart will be successful."` You'll create this Metric-based SLO next.

1. In Datadog, return to **Service Mgmt > Services > SLOs** and click **New SLO** in the upper-right.
2. Under **Define your SLO measurement**, select **By Count**. Below this, notice the **Define your SLI** heading with two sets of form fields: one for **Good events** and another for **Total events**.

1

Define your SLO measurement

Select how to measure your SLO

✓ By Count

Measures the ratio of good events to total events.

🕒 By Monitor Uptime

Measures the uptime of your monitors.

Define your SLI

Count-based SLOs are calculated using a percentage of good events out of total events (Good Events ÷ Total Events × 100). [Need help? See our docs](#)

Good Events

a

(select metric)

from

(everywhere)

sum by

(everything)

</>

as count

+ Add Query

+ Add Formula

Total Events

a

(select metric)

from

(everywhere)

sum by

(everything)

</>

as count


+ Add Query

+ Add Formula

3. Under **Good events**, in the **a** field, insert `trace.rack.request.hits`. Then update each field as follows:

- In the **from** field, insert each of the following:
 - `env :`
 - `service:your-service`
 - `resource_name`
- Leave **sum by** blank.

Define your SLI

Count-based SLOs are calculated using a percentage of good events out of total events (Good Events ÷ Total Events × 100). [Need help? See our docs](#) 

Good Events

a

trace.rack.request.hits

from

</>

env:ruby-shop × , service:store-frontend × ,
resource_name:spree::orderscontroller_edit ×


sum by (everything) as count

+ Add Query

+ Add Formula

4. Below the fields you just completed, click **Add Query**.

Define your SLI

Count-based SLOs are calculated using a percentage of good events out of total events (Good Events ÷ Total Events × 100). [Need help? See our docs](#) 

Good Events

a

trace.rack.request.hits

from

</>

env:ruby-shop × , service:store-frontend × ,
resource_name:spree::orderscontroller_edit ×

sum by (everything) as count

+ Add Query

+ Add Formula

5. Additional fields will appear. Complete them as follows:

- For **b**, replace the pre-populated value with `trace.rack.request.errors`.
- In the **from** field, confirm the following are present:
 - `env:`
 - `service:your-service`
 - `resource_name`
- Leave **sum by** blank.

Define your SLI

Count-based SLOs are calculated using a percentage of good events out of total events (Good Events ÷ Total Events × 100). [Need help? See our docs](#)

Good Events

a	trace.rack.request.hits	from	</> x
env:ruby-shop x ,service:store-frontend x , resource_name:spree::orderscontroller_edit x			
sum by		(everything)	as count

b	trace.rack.request.errors	from	</> x
env:ruby-shop x ,service:store-frontend x , resource_name:spree::orderscontroller_edit x			
sum by		(everything)	as count

6. At the bottom of **Good events**, locate the field reading `a + b`. Update this to `a - b`.

Define your SLI

Count-based SLOs are calculated using a percentage of good events out of total events (Good Events ÷ Total Events × 100). [Need help? See our docs](#)

Good Events

a	trace.rack.request.hits	from	</> x
env:ruby-shop x ,service:store-frontend x , resource_name:spree::orderscontroller_edit x			
sum by		(everything)	as count

b	trace.rack.request.errors	from	</> x
env:ruby-shop x ,service:store-frontend x , resource_name:spree::orderscontroller_edit x			
sum by		(everything)	as count

→

a - b

x

+ Add Query

+ Add Formula

Note: There is no direct metric for successful requests. Instead, there are metrics for total requests (`trace.rack.request.hits`) and for requests with errors (`trace.rack.request.errors`). You can calculate the number of successful requests by subtracting error-ridden requests from total requests. That is,

`trace.rack.request.hits - trace.rack.request.errors` . Or, in this context, `a - b` .

7. Under **Total events**, update fields as follows:

- **a:** `trace.rack.request.hits`
- **from:** Add each of the following:
 - `env:`
 - `service:your-service`
 - `resource_name`
- **sum by:** Leave blank.

Total Events

a

trace.rack.request.hits

from

</>

env:ruby-shop ×

, service:store-frontend ×

, resource_name:spree::orderscontroller_edit ×

sum by

(everything)

as count

+ Add Query

+ Add Formula

8. Under **Set your target & time window**, update the statement to read: **Evaluate over a rolling time window of** `30 days` **with a target of** `99%` . Leave the optional **Warning threshold** blank.

Set your target & time window

Evaluate over a rolling time window of

30 Days ▾

with a target of

99

%

OPTIONAL

Warning threshold

(99.95)

%

9. Under **Add name and tags**, enter the following **Name**:

Comparing Requests

For Description, copy/paste the following:

99% of requests will be successful

Select each of the following **Tags**:

- `env:`
- `service:your-service`
- `resource_name`

Add **Team Tags** as-needed.

Add name and tags

Name

Description Edit Preview

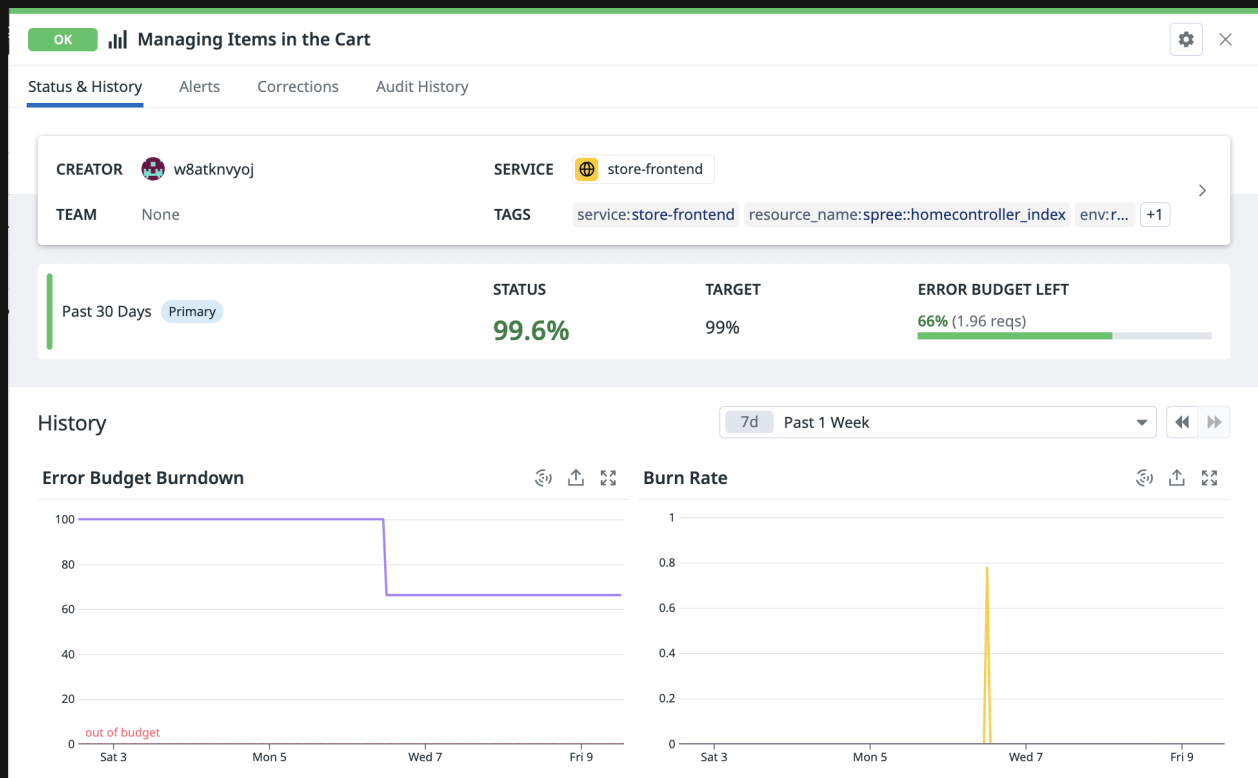
H B I S | | |

99% of requests to the cart will be successful

Tags

Team Tags

Click **Create**. You will be directed to detail panel for the new SLO.

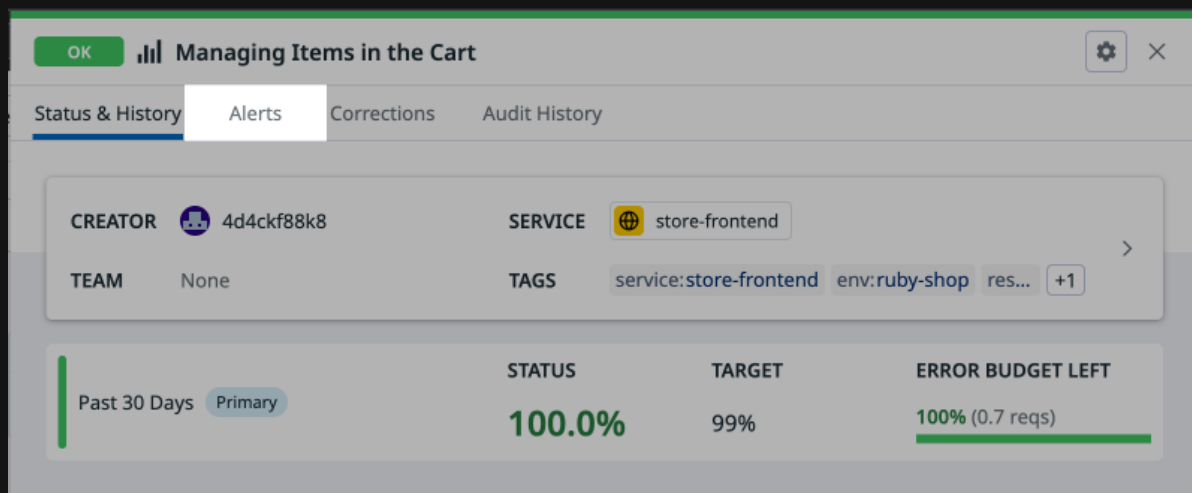


Creating an error budget

Error budget monitors notify you when an SLO's error budget has been consumed to (or past) a point of your choosing.

1. In the detail panel for the **SLO you want an error budget for** you just created, click the **Alerts** tab.

Note: You can access this page by navigating to **Service Mgmt > SLOs** and clicking **SLO you want an error budget for**.



2. In the Alerts tab, click the **New Monitor** button on the right. You will navigate to a form.
3. Under **Select SLO**, confirm your **SLO you want an error budget for** SLO is listed.

Select SLO

Managing Items in the Cart

4. Under **Set alert conditions**, select **Error Budget**.

- Update the **Alert if** line to read: **95 % of budget for 30-day target (primary) is consumed.**
- Update the error budget **Warn if** line to read: **90 % is consumed.**

Set alert conditions

Error Budget

Burn Rate

Get notified when you've consumed too much of your error budget. [?](#)

Alert if % of budget for is consumed.

Warn if % is consumed.

5. Under **Notify your team**, enter the following in the subject field:

Error Budget Alert on SLO: SLO you want an error budget for

6. Add a message to the text body

7. Adjust default values as-needed in **Notify your services and your team members**, **Content displayed in notification** and **Renotification** fields.

In the **Tags** field (*NOT* the **Add SLO Tags** option), select each of the following: - `env:{yourenv}`
- `service:your-service` - `resource_name` Leave remaining fields as-is. The completed form should look like similar this:

1 Select SLO

Managing Items in the Cart



[View SLO](#)

2 Set alert conditions

Error Budget

Burn Rate

Get notified when you've consumed too much of your error budget. [?](#)

Alert if % of budget for is consumed.

Warn if % is consumed.

3 Notify your team

Edit

Preview

[? Use Message Template Variables](#)

Error Budget Alert on SLO: Managing Items in the Cart

H B I S | [Link](#) [Quote](#) [Code](#) [List](#) [Checklist](#) [Image](#) [Table](#) [Text](#) [Code](#) [Table](#) [Code](#)

Warn on 90% of error budget consumed. Alert on 95% of error budget consumed.

Notify your services and your team members

Content displayed in notification

Renotification

☐ If this monitor stays in renotify every [?](#)

Tags:

Teams:

Add SLO Tags

Priority:

4 Define permissions and audit notifications

[🔒](#) Restrict editing of this monitor to and [ve264wvnbp@ddtraining.datadoghq.com](#) (creator) [?](#)

If this monitor is modified, notify monitor creator and alert recipients. [?](#)

Test Notifications

Export Monitor

Create

Click **Create**. You will navigate to a detail page for your new monitor. The status will initially read **NO DATA** :

NO DATA Monitor status since 2 seconds ago (9 Jun, 13:29:35 America/Los_Angeles)

Error Budget Alert on SLO: Managing Items in the Cart

▼ Properties



SLO Monitor
ID: 121492484
Created at Jun 9, 2023, 1:29 pm
by w8atknvyoj

TAGS

env:ruby-shop
resource_name:spree::ho...
service:store-frontend

TEAMS

PRIORITY Not Defined

SLO

[Managing Items in the Cart](#)

QUERY

Alert if **100%** of the error budget for the 30-day target is consumed.

```
error_budget("73956a02d994548383513e41bce54544").over("30d") > 100
```

MESSAGE

Warn on 90% of error budget consumed. Alert on 95% of error budget consumed.

It will take several minutes for the monitor to evaluate data and update the status.